

# FraudShieldAI: Insurance Fraud Detection

## 1. Problem Definition

The model aims to detect fraudulent insurance claims based on the provided customer and incident details. It is designed to help insurance companies reduce fraud and prevent potential losses by identifying suspicious claims.

## 2. Data Overview

The model is built using data that includes various features related to the customer's insurance details and the incident for which a claim is made. Some of the important features include:

- **Customer Information:** Age, months as a customer, umbrella limit, policy details (state, CSL, deductible), insured sex and education level, etc.
- **Incident Information:** Type of incident, collision type, incident severity, authorities contacted, property damage, bodily injuries, etc.

The dataset used to train this model contains both categorical (e.g., policy state, incident type) and numerical (e.g., age, months as a customer) features.

## 3. Data Preprocessing

The data preprocessing steps included:

- **Handling Missing Values:** Any missing or invalid values were handled either by imputation or removal.
- **Categorical Encoding:** Categorical variables were one-hot encoded using `pd.get_dummies` to convert them into numerical format.
- **Scaling:** Numerical features were scaled using a standard scaler to ensure that all features are on the same scale.
- **Feature Engineering:** New features were created if necessary, and irrelevant or redundant features were removed.

## 4. Model Selection

For this problem, a classification model is appropriate since the goal is to predict whether a given insurance claim is fraudulent or not. After evaluating several models (e.g., Logistic Regression, Decision Trees, Random Forests), a model was chosen based on performance, interpretability, and computational efficiency.

- **Model Type:** Random Forest (or another classification algorithm, e.g., Gradient Boosting)
- **Reason for Selection:** Random Forests are robust to overfitting, can handle both categorical and numerical features, and work well with imbalanced datasets.

5. Model Training

- **Training Data:** The model was trained on a dataset of insurance claims, with each record labeled as either fraudulent or non-fraudulent.
- **Cross-Validation:** Cross-validation techniques were employed to prevent overfitting and to ensure that the model generalizes well to unseen data.

6. Model Evaluation Metrics

After training the model, the following evaluation metrics were used to assess its performance:

- **Accuracy:** Measures the percentage of correctly classified instances.
- **Precision:** Indicates the proportion of true positive predictions (fraudulent claims correctly identified).
- **Recall (Sensitivity):** Measures how many fraudulent claims were correctly identified by the model.
- **F1-Score:** The harmonic mean of precision and recall, which is useful for imbalanced datasets.
- **AUC-ROC Curve:** A plot to evaluate the tradeoff between sensitivity and specificity across different thresholds.

7. Model Performance

Here are some key performance metrics for the model:

Gradient Boosting Accuracy: 0.74					
Classification Report for Gradient Boosting:					
	precision	recall	f1-score	support	
0	0.80	0.86	0.83	145	
1	0.55	0.44	0.48	55	
accuracy			0.74	200	
macro avg	0.67	0.65	0.66	200	
weighted avg	0.73	0.74	0.74	200	

- **Interpretation:** The model performs well, with a high recall value indicating that most fraudulent claims are correctly identified. The precision value is also reasonably good, showing that the majority of flagged claims are indeed fraudulent. The F1-score suggests a balanced performance across precision and recall.

8. Feature Importance

Feature importance helps understand which features contribute the most to the model's predictions. Some of the most important features include:

- **Policy State:** The state where the policy was issued.
- **Incident Type:** The type of accident (e.g., single-vehicle collision, multi-vehicle collision).
- **Total Claim Amount:** The total amount requested for the claim.
- **Property Damage:** Whether the incident caused property damage.

## 9. Model Deployment

The model has been deployed using Flask, where it is integrated into a web application for real-time fraud prediction. The user inputs the details of an insurance claim, and the model predicts whether the claim is fraudulent or not.

- **Input Form:** The user provides various features like the number of vehicles involved, incident type, policy state, etc.
- **Prediction Output:** The model outputs whether the claim is fraudulent or not, along with the probability of fraud.

## 10. Model Limitations

- **Imbalanced Data:** Insurance fraud detection can be highly imbalanced (with far more legitimate claims than fraudulent ones). The model was trained using techniques like resampling or balancing to mitigate this.
- **External Factors:** The model only uses historical data and may not account for factors that could influence fraud in the future (e.g., new fraud schemes).
- **Data Privacy:** Some of the features used (e.g., customer details, incident information) require ensuring that data privacy regulations are followed.

## 11. Future Improvements

- **Hyperparameter Tuning:** Further tuning of the Random Forest model could help improve performance.
- **Ensemble Methods:** Combining multiple models (e.g., stacking) could provide a better prediction.
- **Feature Engineering:** Adding more relevant features, such as external risk factors, could improve prediction accuracy.
- **Data Augmentation:** Using synthetic data or generating more fraudulent claims could help mitigate class imbalance.

## 12. Conclusion

The insurance fraud detection model performs well in predicting fraudulent claims with a high recall and precision. It can be used in production environments to help insurers detect potential fraud in real-time. The model's deployment as a web application makes it easy for users to interact with and get predictions for individual insurance claims.